

Detroit Blight Ticket Compliance

Goal

- Predict whether a blight ticket will be paid on time.

Steps

- Data loading
- Exploratory Data Analysis
- Baseline model
- Evaluation

1. Import, data loading + Merge

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.metrics import roc_curve, roc_auc_score, classification_report

addresses = pd.read_csv(
    "/Users/User/Desktop/PYTHONPRO/DATA_SCIENCE_KURZUS_4_blokk/notebooks/project_hazi/addresses.csv")
addresses.head()
```

	ticket_id	address
0	22056	2900 tyler, Detroit MI
1	27586	4311 central, Detroit MI
2	22062	1449 longfellow, Detroit MI
3	22084	1441 longfellow, Detroit MI
4	22093	2449 churchill, Detroit MI

```
latlons = pd.read_csv(
    "/Users/User/Desktop/PYTHONPRO/DATA_SCIENCE_KURZUS_4_blokk/notebooks/project_hazi/latlons.csv")
latlons.head()
```

	address	lat	lon
0	4300 rosa parks blvd, Detroit MI 48208	42.346169	-83.079962
1	14512 sussex, Detroit MI	42.394657	-83.194265
2	3456 garland, Detroit MI	42.373779	-82.986228
3	5787 wayburn, Detroit MI	42.403342	-82.957805
4	5766 haverhill, Detroit MI	42.407255	-82.946295

```
train = pd.read_csv(
    "/Users/User/Desktop/PYTHONPRO/DATA_SCIENCE_KURZUS_4_blokk/notebooks/project_hazi/train.csv")
```

C:\Users\User\AppData\Local\Temp\ipykernel_36392\3855574826.py:1: DtypeWarning: Columns (11,12,31) have mixed types. Specify dtype option on import or set low_memory=False.
train = pd.read_csv(

```
#=====
# train.csv file violation_street_number transformation original 2900.0 -> new 2900
#=====

train['violation_street_number'] = (
    train['violation_street_number']
    .astype(str)
    .apply(lambda x: str(int(float(x))) if x.replace('.', '').isdigit() else ""))
)
train.head()
```

	ticket_id	agency_name	inspector_name	violator_name	violation_street_number	violation_street_name	violation_zip_code
0	22056	Buildings, Safety Engineering & Env Department	Sims, Martinzie	INVESTMENT INC., MIDWEST MORTGAGE	2900	TYLER	NaN

	ticket_id	agency_name	inspector_name	violator_name	Violation_Street_Number	Violation_Street_Name	Violation_Zip_Code
1	27586	Buildings, Safety Engineering & Env Department	Williams, Darrin	Michigan, Covenant House	4311	CENTRAL	NaN
2	22062	Buildings, Safety Engineering & Env Department	Sims, Martinzie	SANDERS, DERRON	1449	LONGFELLOW	NaN
3	22084	Buildings, Safety Engineering & Env Department	Sims, Martinzie	MOROSI, MIKE	1441	LONGFELLOW	NaN
4	22093	Buildings, Safety Engineering & Env Department	Sims, Martinzie	NATHANIEL, NEAL	2449	CHURCHILL	NaN

5 rows × 34 columns

```
#=====
# train.csv file vilation_street_name lowercase and concatenate
#=====

train['violation_street_name'] = train['violation_street_name'].str.lower().str.strip()

train['address'] = (
    train['violation_street_number'].astype(str) + ' ' +
    train['violation_street_name'].astype(str) + ', Detroit MI'
)

train.head()
```

	ticket_id	agency_name	inspector_name	violator_name	Violation_Street_Number	Violation_Street_Name	Violation_Zip_Code
0	22056	Buildings, Safety Engineering & Env Department	Sims, Martinzie	INVESTMENT INC., MIDWEST MORTGAGE	2900	tyler	NaN
1	27586	Buildings, Safety Engineering & Env Department	Williams, Darrin	Michigan, Covenant House	4311	central	NaN
2	22062	Buildings, Safety Engineering & Env Department	Sims, Martinzie	SANDERS, DERRON	1449	longfellow	NaN
3	22084	Buildings, Safety Engineering & Env Department	Sims, Martinzie	MOROSI, MIKE	1441	longfellow	NaN
4	22093	Buildings, Safety Engineering & Env Department	Sims, Martinzie	NATHANIEL, NEAL	2449	churchill	NaN

5 rows × 35 columns

```
#=====
# train.csv és latlons.csv merge
#=====

df = train.merge(latlons, on='address', how='left')
df.head()
```

	ticket_id	agency_name	inspector_name	violator_name	violation_street_number	violation_street_name	violation_zip_code
0	22056	Buildings, Safety Engineering & Env Department	Sims, Martinzie	INVESTMENT INC., MIDWEST MORTGAGE	2900	tyler	NaN
1	27586	Buildings, Safety Engineering & Env Department	Williams, Darrin	Michigan, Covenant House	4311	central	NaN
2	22062	Buildings, Safety Engineering & Env Department	Sims, Martinzie	SANDERS, DERRON	1449	longfellow	NaN
3	22084	Buildings, Safety Engineering & Env Department	Sims, Martinzie	MOROSI, MIKE	1441	longfellow	NaN
4	22093	Buildings, Safety Engineering & Env Department	Sims, Martinzie	NATHANIEL, NEAL	2449	churchill	NaN

5 rows × 37 columns



2. EDA + missing values analysis

```
df.isnull().sum()
```

ticket_id	0
agency_name	0
inspector_name	0
violator_name	34
violation_street_number	0
violation_street_name	0
violation_zip_code	250306
mailing_address_str_number	3602
mailing_address_str_name	4
city	0
state	93
zip_code	1
non_us_str_code	250303
country	0
ticket_issued_date	0
hearing_date	12491
violation_code	0
violation_description	0
disposition	0
fine_amount	1
admin_fee	0
state_fee	0
late_fee	0
discount_amount	0
clean_up_cost	0
judgment_amount	0
payment_amount	0
balance_due	0
payment_date	209193
payment_status	0
collection_status	213409
grafitti_status	250305
compliance_detail	0
compliance	90426
address	0
lat	3
lon	3

dtype: int64

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250306 entries, 0 to 250305
Data columns (total 37 columns):
```

#	Column	Non-Null Count	Dtype
0	ticket_id	250306	non-null int64
1	agency_name	250306	non-null object
2	inspector_name	250306	non-null object
3	violator_name	250272	non-null object
4	violation_street_number	250306	non-null object
5	violation_street_name	250306	non-null object
6	violation_zip_code	0	non-null float64
7	mailing_address_str_number	246704	non-null float64
8	mailing_address_str_name	250302	non-null object
9	city	250306	non-null object
10	state	250213	non-null object
11	zip_code	250305	non-null object
12	non_us_str_code	3	non-null object
13	country	250306	non-null object
14	ticket_issued_date	250306	non-null object
15	hearing_date	237815	non-null object
16	violation_code	250306	non-null object
17	violation_description	250306	non-null object
18	disposition	250306	non-null object
19	fine_amount	250305	non-null float64
20	admin_fee	250306	non-null float64
21	state_fee	250306	non-null float64
22	late_fee	250306	non-null float64
23	discount_amount	250306	non-null float64
24	clean_up_cost	250306	non-null float64
25	judgment_amount	250306	non-null float64
26	payment_amount	250306	non-null float64
27	balance_due	250306	non-null float64
28	payment_date	41113	non-null object
29	payment_status	250306	non-null object
30	collection_status	36897	non-null object
31	grafitti_status	1	non-null object
32	compliance_detail	250306	non-null object
33	compliance	159880	non-null float64
34	address	250306	non-null object
35	lat	250303	non-null float64
36	lon	250303	non-null float64

dtypes: float64(14), int64(1), object(22)

memory usage: 70.7+ MB

```
#=====
# Features Nan values summary
=====

df.isna().sum().sort_values(ascending=False)
```

violation_zip_code	250306
grafitti_status	250305
non_us_str_code	250303
collection_status	213409
payment_date	209193
compliance	90426
hearing_date	12491
mailing_address_str_number	3602
state	93
violator_name	34
mailing_address_str_name	4
lon	3
lat	3
fine_amount	1
zip_code	1
agency_name	0
ticket_id	0
violation_street_number	0
inspector_name	0
city	0
violation_street_name	0
admin_fee	0
disposition	0
violation_description	0
violation_code	0
country	0
ticket_issued_date	0
discount_amount	0
state_fee	0
balance_due	0
payment_amount	0
judgment_amount	0
clean_up_cost	0
late_fee	0
payment_status	0
compliance_detail	0

address
dtype: int64

```
#=====
# Leakage és irreleváns oszlopok (jellemzők) eldobása
#=====

leak_cols = [
    'payment_date',
    'payment_amount',
    'payment_status',
    'balance_due',
    'compliance_detail'
]

df = df.drop(columns=leak_cols, errors='ignore')

drop_cols = [
    'violation_zip_code',
    'grafitti_status',
    'non_us_str_code',
    'collection_status'
]

df = df.drop(columns=drop_cols)
df.head()
```

	ticket_id	agency_name	inspector_name	violator_name	violation_street_number	violation_street_name	mailing_address_s
0	22056	Buildings, Safety Engineering & Env Department	Sims, Martinzie	INVESTMENT INC., MIDWEST MORTGAGE	2900	tyler	3.0
1	27586	Buildings, Safety Engineering & Env Department	Williams, Darrin	Michigan, Covenant House	4311	central	2959.0
2	22062	Buildings, Safety Engineering & Env Department	Sims, Martinzie	SANDERS, DERRON	1449	longfellow	23658.0
3	22084	Buildings, Safety Engineering & Env Department	Sims, Martinzie	MOROSI, MIKE	1441	longfellow	5.0
4	22093	Buildings, Safety Engineering & Env Department	Sims, Martinzie	NATHANIEL, NEAL	2449	churchill	7449.0

5 rows × 28 columns

◀

```
df.isnull().sum()
```

ticket_id	0
agency_name	0
inspector_name	0
violator_name	34
violation_street_number	0
violation_street_name	0
mailing_address_str_number	3602
mailing_address_str_name	4
city	0
state	93
zip_code	1
country	0
ticket_issued_date	0
hearing_date	12491
violation_code	0

```
violation_description      0
disposition                0
fine_amount                 1
admin_fee                  0
state_fee                  0
late_fee                   0
discount_amount             0
clean_up_cost               0
judgment_amount             0
compliance                  90426
address                     0
lat                         3
lon                         3
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250306 entries, 0 to 250305
Data columns (total 28 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ticket_id        250306 non-null   int64  
 1   agency_name      250306 non-null   object  
 2   inspector_name   250306 non-null   object  
 3   violator_name    250272 non-null   object  
 4   violation_street_number  250306 non-null   object  
 5   violation_street_name  250306 non-null   object  
 6   mailing_address_str_number  246704 non-null   float64 
 7   mailing_address_str_name   250302 non-null   object  
 8   city              250306 non-null   object  
 9   state             250213 non-null   object  
 10  zip_code          250305 non-null   object  
 11  country           250306 non-null   object  
 12  ticket_issued_date  250306 non-null   object  
 13  hearing_date     237815 non-null   object  
 14  violation_code    250306 non-null   object  
 15  violation_description  250306 non-null   object  
 16  disposition       250306 non-null   object  
 17  fine_amount       250305 non-null   float64 
 18  admin_fee          250306 non-null   float64 
 19  state_fee          250306 non-null   float64 
 20  late_fee           250306 non-null   float64 
 21  discount_amount    250306 non-null   float64 
 22  clean_up_cost      250306 non-null   float64 
 23  judgment_amount    250306 non-null   float64 
 24  compliance          159880 non-null   float64 
 25  address            250306 non-null   object  
 26  lat                250303 non-null   float64 
 27  lon                250303 non-null   float64 
dtypes: float64(11), int64(1), object(16)
memory usage: 53.5+ MB
```

3. Feature engineering

```
#=====
# Target creation
#=====

df = df[df['compliance'].isin([0,1])].copy()
df['target'] = df['compliance']
df.drop(columns=['compliance'], inplace=True)

#=====
# Datums transformation
#=====

df['ticket_issued_date'] = pd.to_datetime(df['ticket_issued_date'])
df['hearing_date'] = pd.to_datetime(df['hearing_date'], errors='coerce')

df['days_until_hearing'] = (
    df['hearing_date'] - df['ticket_issued_date']
).dt.days.fillna(-1)

df['issue_year'] = df['ticket_issued_date'].dt.year
df['issue_month'] = df['ticket_issued_date'].dt.month
df['issue_dow'] = df['ticket_issued_date'].dt.dayofweek

#=====
# Log-transzformation
#=====

money_cols = ['fine_amount', 'admin_fee', 'state_fee', 'late_fee', 'judgment_amount']
for c in money_cols:
    if c in df.columns:
        df[f"log_{c}"] = np.log1p(df[c].fillna(0))

#=====
# Address / geo feature engineering
#=====

# lat (latitude) column checking, if it there are coordinates ->1 otherwise ->0
df['has_latlon'] = df['lat'].notna().astype(int)

df['prior_tickets_by_address'] = df.groupby(
    ['violation_street_name', 'violation_street_number']
)['ticket_id'].transform('count')

df['prior_tickets_by_violator'] = df.groupby(
    'violator_name'
)['ticket_id'].transform('count')

#=====
# Feature list
#=====

numeric_feats = [
    'days_until_hearing', 'prior_tickets_by_address',
    'prior_tickets_by_violator', 'has_latlon'
] + [f"log_{c}" for c in money_cols if f"log_{c}" in df.columns]

cat_feats = [
    'violation_code', 'violation_description',
    'disposition', 'agency_name', 'city', 'state'
]
```

4. Baseline Model - Preprocessing & Logistic Regression

```
numeric_transformer = Pipeline([
    ('impute', SimpleImputer(strategy='median')),
    ('scale', StandardScaler())
])

categorical_transformer = Pipeline([
    ('impute', SimpleImputer(strategy='most_frequent')),
    ('ohe', OneHotEncoder(handle_unknown="ignore"))
])

preprocess = ColumnTransformer([
    ('num', numeric_transformer, numeric_feats),
    ('cat', categorical_transformer, cat_feats)
])
```

5. Vizualization


```
#=====
# Korrelációs mátrix
#=====

import seaborn as sns

num_for_corr = [
    'target',
    'fine_amount',
    'judgment_amount',
    'admin_fee',
    'state_fee',
    'prior_tickets_by_address',
    'prior_tickets_by_violator',
    'lat',
    'lon'
]

corr = df[num_for_corr].corr()

plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, cmap='coolwarm', center=0, fmt='.2f', linewidths=0.5)
plt.title('Numerikus változók korrelációs mátrixa', fontsize=14)
plt.show()
```



```
#=====
# Pivot tábla
# - violation_code TOP10
# - compliance_rate átlaggal
#=====

vc_compliance = (
    df.groupby('violation_code')['target']
    .mean()
    .reset_index(name='compliance_rate')
    .sort_values(by= 'compliance_rate', ascending=False)
)

data = vc_compliance.head(10).copy()


# Rövidített címek
data['violation_code_short'] = data['violation_code'].str[:8]
```

```

plt.figure(figsize=(10,6))
sns.barplot(
    data = data,
    x = 'violation_code_short',
    y = 'compliance_rate',
    color ='steelblue'
)

plt.title('Top 10 Violation Code - Highest Compliance Rate')
plt.ylabel('Compliance Rate')
plt.xticks(rotation=45)
plt.ylim(0, 1)
plt.tight_layout()
plt.show()

```



```

# Több numerikus feature aggregálása violation code szerint
pca_features = df.groupby('violation_code').agg({
    'target': 'mean',
    'fine_amount': 'mean',
    'judgment_amount': 'mean',
    'admin_fee': 'mean',
    'state_fee': 'mean'
})

# Hiányzó értékek kitöltése
pca_features = pca_features.fillna(pca_features.mean())
pca_features.head()

```

	target	fine_amount	judgment_amount	admin_fee	state_fee
violation_code					
19420901	0.250000	475.000000	549.375000	20.0	10.0
19450901	0.027778	309.753086	369.950617	20.0	10.0
19830901	0.000000	47.500000	82.000000	20.0	10.0
19840901	0.500000	25.000000	57.500000	20.0	10.0
19850901	0.000000	50.000000	85.000000	20.0	10.0

```

#=====
# Scaler before PCA
#=====

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
scaled = scaler.fit_transform(pca_features)

```

```

#=====
# PCA futtatása
#=====

from sklearn.decomposition import PCA

# Csak 2 komponens kell a vizualizációhoz
pca = PCA(n_components=2)
pca_result = pca.fit_transform(scaled)

# Hozzáadjuk az eredményt a dataframe-hez
pca_features['PC1'] = pca_result[:,0]
pca_features['PC2'] = pca_result[:,1]

pca_features.head()

```


	target	fine_amount	judgment_amount	admin_fee	state_fee	PC1	PC2
violation_code							
19420901	0.250000	475.000000	549.375000	20.0	10.0	0.267544	1.327232
19450901	0.027778	309.753086	369.950617	20.0	10.0	-0.064850	-0.376539
19830901	0.000000	47.500000	82.000000	20.0	10.0	-0.496860	-0.574262
19840901	0.500000	25.000000	57.500000	20.0	10.0	-0.382449	3.283952
19850901	0.000000	50.000000	85.000000	20.0	10.0	-0.492634	-0.574428

```

import matplotlib.pyplot as plt

plt.figure(figsize=(7, 5))
plt.plot(
    range(1, len(pca.explained_variance_ratio_)+1),
    pca.explained_variance_ratio_,
    marker='o'
)
plt.title('Scree Plot - PCA Variance Explained')
plt.xlabel('Principal Component')
plt.ylabel('Proportion of Variance Explained')
plt.grid(alpha=0.3)
plt.show()


```



```
plt.figure(figsize=(10, 7))

plt.scatter(
    pca_features['PC1'],
    pca_features['PC2'],
    c=pca_features["target"],
    cmap='viridis',
    alpha=0.9
)

plt.title('Violation Code - PCA 2D Projection')
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.colorbar(label='Compliance Rate')
plt.grid(alpha=0.3)
plt.show()
```



PCA Scree Plot

Shows the proportion of variance explained by each principal component, helping us understand how many components capture most of the information.

2D PCA Projection

Visualizes the first two principal components colored by target variable (Compliance).
This highlights clusters and separation patterns in the data.

=====

6. Modells

```
#=====
# Függvény a modell betanításhoz
# Train model on train/test split and return:
# - probabilities indexed by ticket_id
# - AUC score
#=====

from sklearn.model_selection import train_test_split


def train_and_predict(df, model):
    X = df.drop(columns=['target'])
    y = df['target']

    # ticket_id külön kezelése visszaadás miatt
    ticket_ids = X['ticket_id'].copy()
    X = X.drop(columns=['ticket_id'])

    X_train, X_test, y_train, y_test, id_train, id_test = train_test_split(
        X, y, ticket_ids,
        test_size=0.2, random_state=42, stratify=y
    )

    model.fit(X_train, y_train)
    y_pred_proba = model.predict_proba(X_test)[:,1]

    pred_series = pd.Series(y_pred_proba, index=id_test, name='probability')

    auc = roc_auc_score(y_test, y_pred_proba)

    print(f"Model AUC: {auc:.4f}")

    return pred_series, auc, y_test
```

```
#=====
# Modell1: Logistic Regression (baseline)
#=====

from sklearn.linear_model import LogisticRegression

model_lr = Pipeline([
    ('prep', preprocess),
    ('clf', LogisticRegression(max_iter=200))
])

pred_lr, auc_lr, y_test_lr = train_and_predict(df, model_lr)
pred_lr.head()
```

Model AUC: 0.8139

```
ticket_id
138138    0.035212
149161    0.022452
57262     0.357084
30505     0.061585
192339    0.013783
Name: probability, dtype: float64
```

```
#=====
# Modell 2: RandomForest
#=====

from sklearn.ensemble import RandomForestClassifier

model_rf = Pipeline([
    ("prep", preprocess),
    ('clf', RandomForestClassifier(
        n_estimators=100,
        random_state=42,
        max_depth=15))
])

pred_rf, auc_rf, y_test_rf = train_and_predict(df, model_rf)
pred_rf.head()
```

Model AUC: 0.8171

```
ticket_id
138138    0.046453
149161    0.046118
57262     0.203390
30505     0.066659
192339    0.035467
Name: probability, dtype: float64
```

```
=====
# Modell 3: GradientBoostingClassifier
=====

from sklearn.ensemble import GradientBoostingClassifier

model_gb = Pipeline([
    ("prep", preprocess),
    ('clf', GradientBoostingClassifier())
])

pred_gb, auc_gb, y_test_gb = train_and_predict(df, model_gb)
pred_rf.head()
```

```
Model AUC: 0.8314
```

```
ticket_id
138138    0.046453
149161    0.046118
57262     0.203390
30505     0.066659
192339    0.035467
Name: probability, dtype: float64
```

Interpretation

Gradient Boosting achieved the highest ROC-AUC score, indicating better ranking performance compared to Logistic Regression and Random Forest.

Given the class imbalance, ROC-AUC is a more appropriate metric than accuracy.

7. Evaluation

```
results = pd.DataFrame({
    'Model': ['Logistic Regression', 'Random Forest', 'Gradient Boosting'],
    'ROC AUC': [auc_lr, auc_rf, auc_gb]
})

results.style.background_gradient(cmap='Blues').format({'AUC': '{:.4f}'})
```

Model	ROC AUC
0 Logistic Regression	0.813936
1 Random Forest	0.817129
2 Gradient Boosting	0.831400

```
import matplotlib.pyplot as plt


def plot_roc_curve(y_true, y_pred_proba, model_name='Model'):
    fpr, tpr, thresholds = roc_curve(y_true, y_pred_proba)
    auc = roc_auc_score(y_true, y_pred_proba)

    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, label=f'{model_name} AUC = {auc:.4f}')
    plt.plot([0, 1], [0, 1], linestyle='--', label='Random baseline')
    plt.xlabel('False positive rate')
    plt.ylabel('True positive rate')
    plt.title(f'ROC Curve - {model_name}')
    plt.legend()
    plt.grid(alpha=0.3)
    plt.show()
```

```
pred_lr, auc_lr, y_test_lr = train_and_predict(df, model_lr)

plot_roc_curve(y_test_lr, pred_lr.values, "LogisticRegression")
```

Model AUC: 0.8139



Classification reports for Logistic Regression:

```
def print_classification_report(y_test, probas, threshold=0.5):
    y_pred = (probas >= threshold).astype(int)
    print("Classification report for Logistic Regression:\n")
    print(classification_report(y_test, y_pred))

print_classification_report(y_test_lr, pred_lr.values)
```

Classification report for Logistic Regression:


precision	recall	f1-score	support
-----------	--------	----------	---------

0.0	0.94	1.00	0.97	29657
1.0	0.82	0.15	0.26	2319
accuracy			0.94	31976
macro avg	0.88	0.58	0.61	31976
weighted avg	0.93	0.94	0.92	31976

```
pred_rf, auc_rf, y_test_rf =train_and_predict(df, model_rf)

plot_roc_curve(y_test_rf, pred_rf.values, "Random Forest")
```

Model AUC: 0.8171



Classification reports for Random Forest:

```
def print_classification_report(y_test, probas, threshold=0.5):
    y_pred = (probas >= threshold).astype(int)
    print("Classification report for Random Forest:\n")
    print(classification_report(y_test, y_pred))

print_classification_report(y_test_rf, pred_rf.values)
```


Classification report for Random Forest:

	precision	recall	f1-score	support
0.0	0.93	1.00	0.96	29657
1.0	1.00	0.00	0.01	2319
accuracy			0.93	31976
macro avg	0.96	0.50	0.49	31976
weighted avg	0.93	0.93	0.89	31976

```
pred_gb, auc_gb, y_test_gb =train_and_predict(df, model_gb)

plot_roc_curve(y_test_gb, pred_gb.values, "Gradient Boosting")
```

Model AUC: 0.8314



Classification reports for Gradient Boosting:

```
def print_classification_report(y_test, probas, threshold=0.5):
    y_pred = (probas >= threshold).astype(int)
    print("Classification report for Gradient Boosting:\n")
    print(classification_report(y_test, y_pred))

print_classification_report(y_test_gb, pred_gb.values)
```

Classification report for Gradient Boosting:

	precision	recall	f1-score	support
0.0	0.94	1.00	0.97	29657
1.0	0.84	0.22	0.35	2319
accuracy			0.94	31976
macro avg	0.89	0.61	0.66	31976
weighted avg	0.94	0.94	0.92	31976

